

# Technical Interview Exercise: Mining and Structuring Drug Indications from Labels

## Objective

Develop a **microservice-based application** that extracts drug indications from **DailyMed** drug labels, maps them to standardized medical vocabulary (ICD-10 codes), and provides a **queryable API**. The implementation must be in **Python, Node.js, or .NET** and follow **enterprise-grade software principles**, including:

- **Test-Driven Development (TDD)**
- **Clean Architecture** (separation of concerns, layered design)
- **High Code Quality** (readability, maintainability, modularity)
- **Scalability & Performance Considerations**
- **Dockerized Deployment** (with `docker-compose` for execution)

## Requirements

### 1. Core Features

#### Data Extraction

- Scrape or parse **DailyMed** drug labels for **Dupixent**.
  - Extract relevant sections describing **indications**.

#### Indication Processing & Mapping

- Map extracted indications to **ICD-10** codes using an open-source dataset.
- Handle edge cases like:
  - **Synonyms** (e.g., "Hypertension" vs. "High Blood Pressure").
  - **Drugs with multiple indications**.
  - **Unmappable conditions**.
- This should use AI/LLM to complete this step.

#### Structured Data Output

- Store structured drug-indication mappings in a **database or NoSQL store**.
- Make mappings queryable via an API.

## 2. Enterprise-Grade API

- **Develop a Web API** using **.NET (C#)**, **Python (FastAPI/Flask)**, or **Node.js (Express/NestJS)**.
- Implement **CRUD operations**:
  - Create, read, update, and delete drug-indication mappings.
- **Authentication & Authorization**
  - Users should be able to register and log in.
  - Implement role-based access control.
- Include **Swagger or Postman workspace** for API testing.
- Ensure consistent data types (e.g., **true/false**, numbers as strings).
- Implement validation rules for missing or ambiguous data.
- Provides an endpoint (**/programs/<program\_id>**) returning structured JSON.
- Supports querying program details dynamically.

## 3. Data & Storage Layer

- Use a **database** (SQL or NoSQL) to store:
  - Drug-indication mappings
  - User authentication data

## 4. Business Logic Layer

- Keep **business rules** independent of the API and data layers.
- Implement validation logic for incoming data.

## 5. Testing & Quality

- Follow **TDD**: write unit tests **before** implementation.
- Cover:
  - Data extraction and processing logic.
  - API endpoints.
  - Business rules.
  - Authentication flows.
- Ensure **high test coverage**.

## Deliverables

1. **GitHub Repository** containing:
  - **Source code** for the full project.
  - **Unit tests** for API, business logic, and data handling.
  - **README.md** with detailed setup and execution instructions.
2. **README.md must include:**

- **Step-by-step setup** for running the project.
  - **API documentation.**
  - **Sample output** of the system.
  - **Scalability considerations.**
  - **Potential improvements & production challenges.**
3. **Dockerized Deployment**
- Project must be runnable using `docker-compose up` as the only setup step.
4. **Answer this short prompt:**
- How would you lead an engineering team to implement and maintain this project?

## Evaluation Criteria

Category	Description
Clean Architecture	Separation of concerns, modularity, maintainability.
Test-Driven Development	Unit tests for API, business logic, data layer.
Code Quality	Readability, documentation, adherence to best practices.
Functionality	API correctness, data extraction accuracy, ICD-10 mapping.
Scalability & Design	Consideration for large-scale use, error handling.
Dockerization	Ability to launch project using only <code>docker-compose up</code> .
Presentation (Interview)	Clear walkthrough of code, choices, and trade-offs.

## Bonus Points

- **Data Enrichment:**
  - Implement additional **rules-based logic** for missing fields.
- **Performance Optimization:**
  - Preprocess and **cache structured data** for API efficiency.
  - Allow filtering results dynamically based on parameters.
  - Implement **rate-limiting and security best practices**.
- **Gap Analysis:**
  - What did we not think of?
  - What edge cases did you find?
  - What improvements would you recommend?
  - For each gap:
    - Would you:
      - Adapt the data model?
      - Abstract it at the service or API layer?
    - Why?
    - What are the tradeoffs in performance, maintainability, or team clarity?

## Submission & Interview

1. **Submit your GitHub repo link.**
2. **Prepare a Zoom presentation:**
  - Walk through your **user story**, **architecture**, and **technical decisions**.
  - Demo API functionality.
  - Answer code review questions.

Good luck!